

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
dataset=pd.read_csv("information.csv")
```

```
dataset.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 511 entries, 0 to 510
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CRIM        511 non-null    float64
1   ZN          511 non-null    float64
2   INDUS       511 non-null    float64
3   CHAS        511 non-null    int64
4   NOX         511 non-null    float64
5   RM          511 non-null    float64
6   AGE         511 non-null    float64
7   DIS         511 non-null    float64
8   RAD         511 non-null    int64
9   TAX         511 non-null    int64
10  PTRATIO     511 non-null    float64
11  B           511 non-null    float64
12  LSTAT       511 non-null    float64
13  MEDV        511 non-null    float64
```

```
dtypes: float64(11), int64(3)
memory usage: 56.0 KB
```

```
dataset.describe()
```

	CRIM	ZN	INDUS	CHAS	NOX
RM \					
count	511.000000	511.000000	511.000000	511.000000	511.000000
mean	3.584139	11.252446	11.151096	0.068493	0.554757
std	8.564433	23.234838	6.828175	0.252838	0.115310
min	0.006320	0.000000	0.460000	0.000000	0.385000
25%	0.082325	0.000000	5.190000	0.000000	0.449000
50%	0.261690	0.000000	9.690000	0.000000	0.538000
75%	3.621175	12.500000	18.100000	0.000000	0.624000
max	88.976200	100.000000	27.740000	1.000000	0.871000

	AGE	DIS	RAD	TAX	PTRATIO
B \					
count	511.000000	511.000000	511.000000	511.000000	511.000000
mean	68.616243	3.783876	9.485323	407.440313	18.500000
std	28.099130	2.098631	8.688469	167.903532	2.200348
min	2.900000	1.129600	1.000000	187.000000	12.600000
25%	45.050000	2.100350	4.000000	279.500000	17.400000
50%	77.300000	3.152300	5.000000	330.000000	19.100000
75%	94.050000	5.118000	24.000000	666.000000	20.200000
max	100.000000	12.126500	24.000000	711.000000	23.000000

	LSTAT	MEDV
count	511.000000	511.000000
mean	12.879550	22.682192
std	7.797416	9.484262
min	1.730000	5.000000
25%	7.065000	17.050000
50%	11.450000	21.200000

```
75%      17.105000    25.000000
max       76.000000    67.000000
```

```
dataset.corr()
```

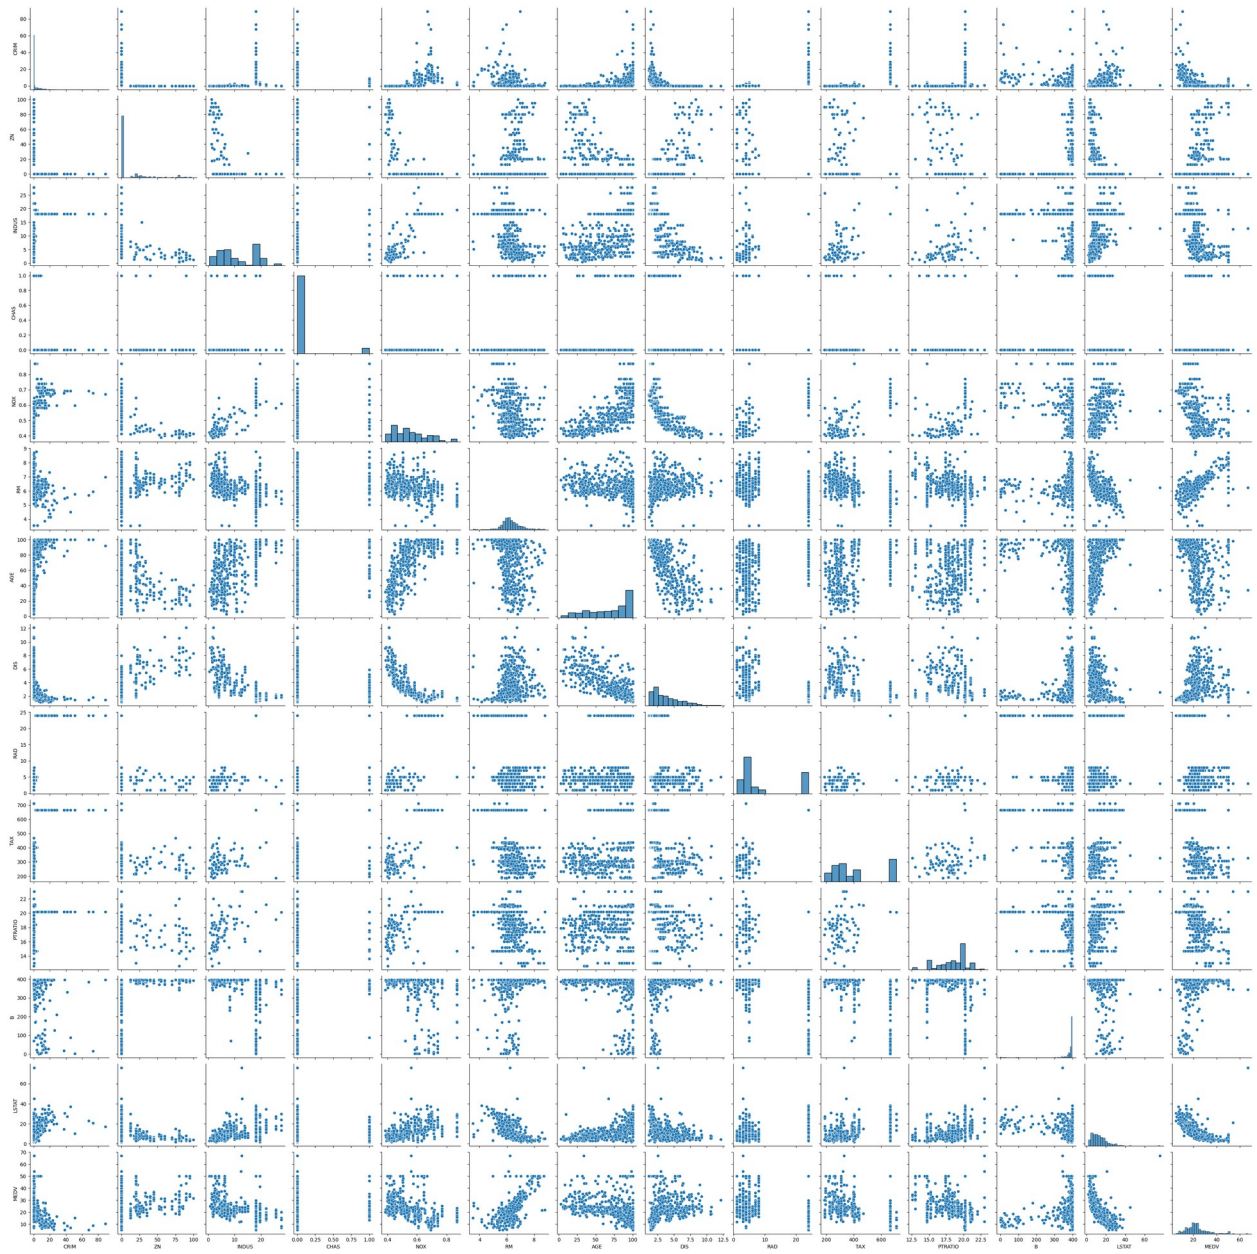
	CRIM	ZN	INDUS	CHAS	NOX	RM
AGE \						
CRIM	1.000000	-0.198451	0.405524	-0.054906	0.420524	-0.206994
0.350847						
ZN	-0.198451	1.000000	-0.534106	-0.041333	-0.516256	0.288855
0.567796						
INDUS	0.405524	-0.534106	1.000000	0.062332	0.763583	-0.370202
0.642817						
CHAS	-0.054906	-0.041333	0.062332	1.000000	0.091023	0.089403
0.085814						
NOX	0.420524	-0.516256	0.763583	0.091023	1.000000	-0.285351
0.729233						
RM	-0.206994	0.288855	-0.370202	0.089403	-0.285351	1.000000
0.227381						
AGE	0.350847	-0.567796	0.642817	0.085814	0.729233	-0.227381
1.000000						
DIS	-0.377028	0.665185	-0.707886	-0.097541	-0.768309	0.166746
0.745097						
RAD	0.625964	-0.307146	0.591784	-0.005343	0.609343	-0.203099
0.452229						
TAX	0.583389	-0.311531	0.718764	-0.034244	0.666982	-0.274950
0.503360						
PTRATIO	0.276695	-0.392844	0.379441	-0.124415	0.186073	-0.320076
0.258251						
B	-0.384356	0.175637	-0.356955	0.048970	-0.379936	0.117844
0.271925						
LSTAT	0.405030	-0.390029	0.556707	-0.057013	0.540050	-0.534628
0.529994						
MEDV	-0.380072	0.339767	-0.463269	0.164782	-0.411486	0.644951
0.368203						

	DIS	RAD	TAX	PTRATIO	B	LSTAT
MEDV						
CRIM	-0.377028	0.625964	0.583389	0.276695	-0.384356	0.405030
0.380072						
ZN	0.665185	-0.307146	-0.311531	-0.392844	0.175637	-0.390029
0.339767						
INDUS	-0.707886	0.591784	0.718764	0.379441	-0.356955	0.556707
0.463269						
CHAS	-0.097541	-0.005343	-0.034244	-0.124415	0.048970	-0.057013
0.164782						
NOX	-0.768309	0.609343	0.666982	0.186073	-0.379936	0.540050
0.411486						
RM	0.166746	-0.203099	-0.274950	-0.320076	0.117844	-0.534628
0.644951						

```
AGE      -0.745097  0.452229  0.503360  0.258251 -0.271925  0.529994 -
0.368203
DIS      1.000000 -0.488474 -0.530379 -0.238155  0.290997 -0.467063
0.233469
RAD      -0.488474  1.000000  0.910211  0.438646 -0.442406  0.422389 -
0.379016
TAX      -0.530379  0.910211  1.000000  0.440962 -0.440830  0.482088 -
0.459274
PTRATIO -0.238155  0.438646  0.440962  1.000000 -0.175251  0.393263 -
0.447464
B        0.290997 -0.442406 -0.440830 -0.175251  1.000000 -0.339910
0.317941
LSTAT    -0.467063  0.422389  0.482088  0.393263 -0.339910  1.000000 -
0.562960
MEDV     0.233469 -0.379016 -0.459274 -0.447464  0.317941 -0.562960
1.000000
```


```
sns.pairplot(df)
```

```
<seaborn.axisgrid.PairGrid at 0x208c256d810>
```




```
sns.heatmap(dataset.corr(),annot = True)
```


```
<Axes: >
```




```
plt.scatter(dataset['CRIM'],dataset['MEDV'])
plt.xlabel("crime rate")
plt.ylabel('medv')
Text(0, 0.5, 'medv')
```




```
plt.scatter(dataset['RM'],dataset['MEDV'])  
<matplotlib.collections.PathCollection at 0x208cf415f90>
```




```
sns.regplot(x="CRIM",y="MEDV",data = dataset)  
<Axes: xlabel='CRIM', ylabel='MEDV'>
```


```
sns.regplot(x="RM",y="MEDV",data = dataset)  
<Axes: xlabel='RM', ylabel='MEDV'>
```



```
sns.boxplot(dataset["MEDV"])  
plt.show()
```



```
sns.boxplot(dataset["CRIM"])\nplt.show()
```



```

def outliers(s):
    q1=df[s].quantile(.25)
    q3=df[s].quantile(.75)
    iqr=q3-q1
    ul=q3+1.5*iqr
    ll=q1-1.5*iqr
    return(ll,ul)

outliers("CRIM")
(np.float64(-5.22595), np.float64(8.92945))

outliers("RM")
(np.float64(4.760749999999999), np.float64(7.750750000000001))

outliers("ZN")
(np.float64(-18.75), np.float64(31.25))

outliers("AGE")
(np.float64(-28.450000000000003), np.float64(167.55))

dataset.shape
(511, 14)

dataset["CRIM"].median()
np.float64(0.26169)

dataset=df[df['CRIM']<8.92945]
dataset

(dataset["CRIM"]>8.92945).sum()
np.int64(67)

X=df.iloc[:, :-1]
y=df.iloc[:, -1]

```

X

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222

4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222
..
506	0.98765	0.0	12.50	0	0.561	6.980	89.0	2.0980	3	320
507	0.23456	0.0	12.50	0	0.561	6.980	76.0	2.6540	3	320
508	0.44433	0.0	12.50	0	0.561	6.123	98.0	2.9870	3	320
509	0.77763	0.0	12.70	0	0.561	6.222	34.0	2.5430	3	329
510	0.65432	0.0	12.80	0	0.561	6.760	67.0	2.9870	3	345

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
..
506	23.0	396.00	12.00
507	23.0	343.00	25.00
508	23.0	343.00	21.00
509	23.0	343.00	76.00
510	23.0	321.00	45.00

[511 rows x 13 columns]

y

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2
..	...
506	12.0
507	32.0
508	54.0
509	67.0
510	24.0

Name: MEDV, Length: 511, dtype: float64

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.33, random_state=10)
```

```
X_train.shape
```

```

(342, 13)
y_train.shape
(342,)
X_test.shape
(169, 13)
y_test.shape
(169,)
dataset["CRIM"]
0      0.00632
1      0.02731
2      0.02729
3      0.03237
4      0.06905
...
506    0.98765
507    0.23456
508    0.44433
509    0.77763
510    0.65432
Name: CRIM, Length: 511, dtype: float64

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()

X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
X_train
array([[ -0.41439929, -0.48202008, -1.26559838, ..., -0.36860961,
         0.38410572,  0.14286711],
       [ -0.16444387, -0.48202008,  1.16969306, ..., -1.76282977,
         0.39184726,  2.19433483],
       [ -0.41792627, -0.48202008, -0.87726044, ...,  0.71078664,
         0.3700401 , -0.95212852],
       ...,
       [ -0.1633557 , -0.48202008,  0.95916553, ...,  0.71078664,
         0.16309014,  0.05843371],
       [ -0.40500056, -0.48202008,  2.03314043, ...,  0.21606336,
         0.31890231,  0.25236667],
       [ -0.34019155,  0.32164046, -1.0508034 , ..., -2.52740211,
         0.39937073, -0.32283585]], shape=(342, 13))

```

X_test

```
array([[ -0.13525904, -0.48202008, -0.20727079, ..., -0.09876055,
        -0.05803447, -0.03391532],
       [ -0.28966984, -0.48202008,  1.16969306, ..., -1.76282977,
         0.08349401, -1.09592915],
       [  0.04254796, -0.48202008,  0.95916553, ...,  0.71078664,
         0.38127079, -0.0299575 ],
       ...,
       [  0.09625216, -0.48202008,  0.95916553, ...,  0.71078664,
        -3.32398389,  0.69036493],
       [-0.42109345,  1.92896154, -1.19874167, ..., -1.35805617,
         0.4099472 , -1.03788119],
       [-0.41238153, -0.48202008,  0.20951682, ...,  0.03616398,
         0.19503763, -0.12230653]], shape=(169, 13))
```

```
from sklearn.linear_model import LinearRegression
```

```
regression=LinearRegression()
```

```
regression.fit(X_train,y_train)
```

```
LinearRegression()
```

```
reg_pred=regression.predict(X_test)
```

reg_pred

```
array([ 19.98589533,  31.45526324,  19.59311181,  24.32210658,
        28.65239095,  12.53813937,  24.07053296,  22.05487354,
        24.89438288,  24.25190669,  32.82469698,  16.34519674,
        21.56455789,  32.37302369,  22.81157581,  20.98905942,
        35.13749789,  36.13891021,  19.06078533,  18.50783589,
        19.61797837,  24.41306731,  19.31342742,  20.63629152,
        27.19826898,  32.17514179,  23.67638675,  21.56756887,
        31.85353367,  35.53316975,  27.66098093,  16.97125538,
       -10.27569514,  25.23276966,  23.36747744,  27.93420839,
        15.46511035,  32.04879176,  14.22574443,  23.29641925,
        24.36276523,  29.27148426,  33.95044511,  21.2986439 ,
        14.8935081 ,  24.79478944,   4.76895317,  19.67801423,
        31.6116077 ,  24.00135581,  19.24845622,  30.34424731,
        20.47177245,  21.33555603,  17.352779 ,  19.39571463,
        23.71026394,  40.07132318,  16.68689291,  28.46186716,
        34.36953395,  34.56498629,  18.15940759,  19.67860881,
        32.69745379,  20.7894055 ,  15.2048199 ,  16.25439973,
        23.7828687 ,  22.56373662,  32.24705347,  39.54029996,
        17.18436899,  15.9986714 ,  19.0917243 ,  21.50077783,
        34.13900845,  30.17485542,  30.1694527 ,  14.43289604,
        37.52295461,  19.13272 ,  24.39150369,  22.92656843,
        13.79039387,  36.06891963,  33.94588756,  27.63690063,
        21.86774329,  19.26172629,  26.37077488,  17.55771661,
```

```

33.93727218, 14.02486744, 23.51597365, 24.69766797,
24.18152769, 19.75581542, 10.27364156, 27.34721329,
29.48391432, 19.06112325, 25.28128478, 17.96120099,
 3.92877823, 21.77087245, 27.8164002 , 15.06357486,
20.9073709 , 30.53538851, 27.01387813, 21.39953587,
23.32641309, 16.9841283 , 23.54058351, 18.30513425,
26.1564794 , 25.07903691, 22.90716366, 9.89417608,
18.26575831, 18.12962688, 34.33410557, 8.67245294,
28.49975845, 24.47791253, 19.99168584, 15.64631398,
16.42932162, 26.0183806 , 20.96568148, 14.57644367,
31.52386014, 26.27641314, 37.9927828 , 19.54774116,
19.74205644, 23.97378827, 24.69811483, 39.65534149,
31.27639108, 27.71310802, 29.17987973, 28.5879573 ,
12.68138443, 18.45594585, 31.23890506, 17.5373893 ,
15.96218637, 14.54511307, 25.95493798, 23.83540554,
27.49694938, 21.89388761, 21.17599861, 34.40904416,
19.87379061, 26.43986821, 27.99812994, 24.88505292,
19.95486094, 15.10219404, 25.65195737, 17.99968273,
19.26950049, 28.12088645, 16.51334319, 32.47883265,
22.74362601])

```

y_test

```

310    16.1
157    41.3
359    22.6
43     24.7
181    36.2
...
460    16.4
308    22.8
455    14.1
193    31.1
76     20.0

```

Name: MEDV, Length: 169, dtype: float64

```

plt.figure(figsize=(10,5))
sns.distplot(reg_pred,hist=False,color="red")
sns.distplot(dataset["MEDV"],hist=False)

```

C:\Users\hp\AppData\Local\Temp\ipykernel_14888\1488617724.py:2:
UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `kdeplot` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>


```
sns.distplot(reg_pred,hist=False,color="red")
C:\Users\hp\AppData\Local\Temp\ipykernel_14888\1488617724.py:3:
UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``kdeplot`` (an axes-level function for kernel density plots).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset["MEDV"],hist=False)
<Axes: xlabel='MEDV', ylabel='Density'>
```



```
plt.scatter(y_test,reg_pred)
plt.xlabel('Test truth Data')
plt.ylabel('Test predicted Data')
Text(0, 0.5, 'Test predicted Data')
```



```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
print(mean_squared_error(y_test,reg_pred))
print(mean_absolute_error(y_test,reg_pred))
print(np.sqrt(mean_squared_error(y_test,reg_pred)))

64.01362373830402
3.992903676554756
8.000851438334799

from sklearn.metrics import r2_score
r2_score(y_test,reg_pred)

0.3065524754881511

from sklearn.linear_model import Ridge
clf = Ridge()
clf.fit(X_train, y_train)
Ridge()
r_pred=clf.predict(X_test)
```

```
y_test
```

```
310    16.1  
157    41.3  
359    22.6  
43     24.7  
181    36.2
```

```
...  
460    16.4  
308    22.8  
455    14.1  
193    31.1  
76     20.0
```

```
Name: MEDV, Length: 169, dtype: float64
```

```
r2_score(y_test, r_pred)
```

```
0.30851570863297906
```