

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import joblib
```

```
pip install joblib
```

Requirement already satisfied: joblib in c:\users\hp\appdata\local\programs\python\python313\lib\site-packages (1.4.2)

Note: you may need to restart the kernel to use updated packages.

```
data = pd.read_csv("bcd.csv")
```

```
data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI \
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	768 non-null	int64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	768 non-null	int64

```
8 Outcome 768 non-null int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
(data.columns)
```

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',
       'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

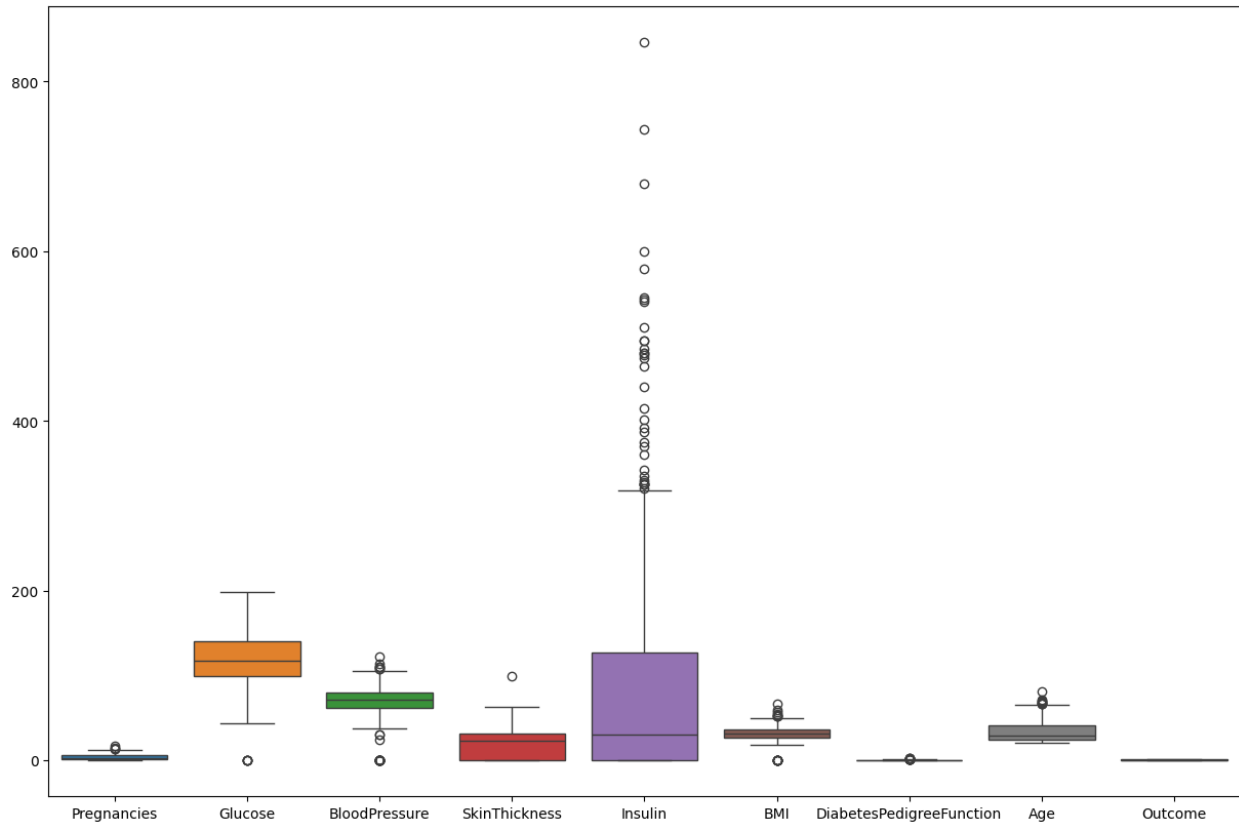
```
data.describe().T
```

	count	mean	std	min
25% \				
Pregnancies	768.0	3.845052	3.369578	0.0000
1.00000				
Glucose	768.0	120.894531	31.972618	0.0000
99.00000				
BloodPressure	768.0	69.105469	19.355807	0.0000
62.00000				
SkinThickness	768.0	20.536458	15.952218	0.0000
0.00000				
Insulin	768.0	79.799479	115.244002	0.0000
0.00000				
BMI	768.0	31.992578	7.884160	0.0000
27.30000				
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.0780
0.24375				
Age	768.0	33.240885	11.760232	21.0000
24.00000				
Outcome	768.0	0.348958	0.476951	0.0000
0.00000				

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
plt.figure(figsize = (15,10))
sns.boxplot(data=data)
```

```
<Axes: >
```



```
# to remove outlayer to find value
def outlier(s):
    q1 = data[s].quantile(.25)
    q3 = data[s].quantile(.75)
    iqr = q3 - q1
    ul = q3+1.5*iqr
    ll = q1-1.5*iqr
    return (ll,ul)

outlier("Glucose")

(np.float64(37.125), np.float64(202.125))

for i in data.columns:
    print(f"{i} = {outlier(i)}")

Pregnancies = (np.float64(-6.5), np.float64(13.5))
Glucose = (np.float64(37.125), np.float64(202.125))
BloodPressure = (np.float64(35.0), np.float64(107.0))
SkinThickness = (np.float64(-48.0), np.float64(80.0))
Insulin = (np.float64(-190.875), np.float64(318.125))
BMI = (np.float64(13.35), np.float64(50.550000000000004))
DiabetesPedigreeFunction = (np.float64(-0.32999999999999999),
np.float64(1.2))
```

```
Age = (np.float64(-1.5), np.float64(66.5))
Outcome = (np.float64(-1.5), np.float64(2.5))
```

```
data.shape
```

```
(768, 9)
```

```
data1 = data[data['Insulin']<280]
```


```
plt.figure(figsize =(10,15))
```

```
<Figure size 1000x1500 with 0 Axes>
```

```
<Figure size 1000x1500 with 0 Axes>
```

```
sns.boxplot(data = data['Insulin'])
```

```
<Axes: ylabel='Insulin'>
```



```
data2 =
data1[(data['Glucose']>40)&((data['BloodPressure']>35)&(data['BloodPressure']<110))]
```


```
C:\Users\hp\AppData\Local\Temp\ipykernel_16728\498802551.py:1:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
data2 =
```

```
data1[(data['Glucose']>40)&((data['BloodPressure']>35)&(data['BloodPressure']<110))]
```

```
plt.figure(figsize =(25,10))
sns.boxplot(data = data2)
```

<Axes: >



```
data.shape
```

```
(768, 9)
```


```
df =
data2[(data2['SkinThickness']<70)&((data['BMI']>17)&(data2['BMI']<50))
&((data['Age']<67)&(data2['Pregnancies']<14))&(data2['DiabetesPedigree
Function']<1.1)]
```

```
C:\Users\hp\AppData\Local\Temp\ipykernel_16728\339148766.py:1:
UserWarning: Boolean Series key will be reindexed to match DataFrame
index.
```

```
df =
data2[(data2['SkinThickness']<70)&((data['BMI']>17)&(data2['BMI']<50))
&((data['Age']<67)&(data2['Pregnancies']<14))&(data2['DiabetesPedigree
Function']<1.1)]
```

```
plt.figure(figsize =(25,10))
sns.boxplot(data = df)
```

<Axes: >



```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 617 entries, 0 to 767

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Pregnancies	617 non-null	int64
1	Glucose	617 non-null	int64
2	BloodPressure	617 non-null	int64
3	SkinThickness	617 non-null	int64
4	Insulin	617 non-null	int64
5	BMI	617 non-null	float64
6	DiabetesPedigreeFunction	617 non-null	float64
7	Age	617 non-null	int64
8	Outcome	617 non-null	int64

```
dtypes: float64(2), int64(7)
```

```
memory usage: 48.2 KB
```

```
x = df[df.columns[:-1]]
```

X

\	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
5	5	116	74	0	0	25.6
...	...	...	...	...	...	...

763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
5	0.201	30
..	...	...
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[617 rows x 8 columns]

df

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
5	5	116	74	0	0	25.6
..	...	...	...	...	...	...
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1

767	1	93	70	31	0	30.4
-----	---	----	----	----	---	------

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
5	0.201	30	0
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[617 rows x 9 columns]

x

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3
3	1	89	66	23	94	28.1
5	5	116	74	0	0	25.6
..	...	...	...	...	...	...
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
5	0.201	30



```

..
763      0.171    63
764      0.340    27
765      0.245    30
766      0.349    47
767      0.315    23

```

```
[617 rows x 8 columns]
```

```

y = df['Outcome']
y

```

```

0      1
1      0
2      1
3      0
5      0

```

```

..
763     0
764     0
765     0
766     1
767     0

```

```
Name: Outcome, Length: 617, dtype: int64
```

```
from sklearn.model_selection import train_test_split
```

```

x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.25,random_state=42)

```

```
x_train.shape
```

```
(462, 8)
```

```
x_test.shape
```

```
(155, 8)
```

```
from sklearn.preprocessing import StandardScaler
```

```
ss = StandardScaler()
```

```

x_train = ss.fit_transform(x_train)
x_train

```

```

array([[ -0.88120816,  0.48512932,  2.56297614, ...,  0.11935477,
        -0.81334446,  1.05677328],
       [ -0.88120816, -1.25338872, -0.36007833, ..., -2.17756899,
         0.88519638, -0.51703163],
       [  1.2566967 ,  2.08456592, -0.01618957, ...,  0.10362241,
         1.29894351,  0.26987083],
       ...,

```

```

[ 0.34045176, 0.10265535, 1.18742109, ..., 0.4182695 ,
 -0.56945142, 0.00757001],
 [-0.88120816, 1.5630105 , 0.84353233, ..., 0.11935477,
 -0.34733454, 1.49394131],
 [-1.18662314, -0.48844078, 1.01547671, ..., -0.6515306 ,
 1.39475863, 2.54314458]], shape=(462, 8))

x_test = ss.transform(x_test)

from sklearn.linear_model import LogisticRegression

classification = LogisticRegression()

classification.fit(x_train,y_train)

LogisticRegression()

cls_pred = classification.predict(x_test)

cls_pred
array([0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
0,
      1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0,
      0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
1,
      0])

y_test
66      1
729     0
104     0
387     1
137     0
..
31      1
366     1
384     0
355     1
312     1
Name: Outcome, Length: 155, dtype: int64

```

```

from sklearn.metrics import
accuracy_score, confusion_matrix, roc_curve, roc_auc_score, ConfusionMatrixDisplay

accuracy_score(y_test, cls_pred)


0.8193548387096774

confusion_matrix(y_test, cls_pred)

array([[104,  1],
       [ 27, 23]])

cm_display =
ConfusionMatrixDisplay(confusion_matrix(y_test, cls_pred)).plot()

```



```

roc_curve(cls_pred, y_test)

(array([0.          , 0.20610687, 1.          ]),
 array([0.          , 0.95833333, 1.          ]),
 array([inf, 1., 0.]))


from sklearn.metrics import RocCurveDisplay, roc_curve
y_score = classification_decision_function(x_test)
fpr, tpr, _ = roc_curve(y_test, y_score, pos_label =

```

```
classification.classes_[1])
roc_display = RocCurveDisplay(fpr = fpr, tpr = tpr).plot()
```

C:\Users\hp\AppData\Local\Programs\Python\Python313\Lib\site-packages\sklearn\metrics\\_plot\roc\_curve.py:189: UserWarning: No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
self.ax_.legend(loc="lower right")
```



```
joblib.dump(classification, "test.pkl")
```

```
['test.pkl']
```

```
xyz = joblib.load('test.pkl')
```

```
xyz
```

```
LogisticRegression()
```

```
joblib.dump(ss, "stand.pkl")
```

```
['stand.pkl']
```